

Optimizing a Minimal Language using Pre-trained Language Models

Ajay Narayanan

Technical University of Munich

Agenda

Introduction

Methodology

System Design and Implementation

Results and Evaluation

Conclusion

Introduction

The Purpose of Language

- Human language evolved to aid survival and propagation.
- Facilitates collaboration, competition, and influence.
- Language is shaped by the trade-off between:
 - **Expressiveness**
 - **Efficiency** (driven by evolution)
- This would lead us to assume natural languages are optimized given these variables.

Why Natural Languages Are Not Optimal

- Real-world communication happens in noisy environments.
- Redundancy helps with error correction.
- Natural languages are also constrained by:
 - Incremental language processing
 - Need for compositionality, systematicity, and concatenation
- Not optimal in the information-theoretic sense.

Constructed Languages (ConLangs)

- Intentionally created languages.
- Types:
 - Auxiliary (e.g. Esperanto)
 - Fictional (e.g. Dothraki, Quenya, Klingon)
 - Logical (e.g. Lojban)
 - Minimalist (e.g. Toki Pona)
 - Expressive (e.g. Ithkuil)

What Does "Optimal" Mean for a Language?

- Information-theoretic efficiency
- Ease of learning
- Expressiveness
- Robustness in noisy conditions
- Ability to convey complex ideas
- **Key Questions:**
 - Can we define and measure these?
 - Can we build languages optimizing for one or more of these constraints?

Can LLMs Help Design Better Languages?

- LLMs encode deep patterns in natural language and world knowledge.
- They may uncover latent linguistic structures.
- Potential use in evaluating or designing:
 - Phonology, orthography, morphology
 - Syntax and vocabulary
- Could LLMs help define and optimize linguistic trade-offs?
- Could LLMs assist in the language generation process?

Thesis Goals

1. Define linguistic optimality across dimensions.
2. Explore design of an efficient constructed language.
3. Investigate how LLMs can aid this process.

Methodology

Methodological Overview

- Step-by-step approach inspired by language construction literature.
- Building a language specification that includes:
 - Phonemic inventory
 - Phonotactics
 - Grammar
 - Vocabulary
- Translation of known texts used for evaluation.

- Modular, computational experimental design.
- Objective: Guide LLMs to generate human-like ConLangs.
- Pipeline consists of:
 - Phonology
 - Morphology
 - Syntax
 - Vocabulary
- Allows ablative analysis and future extensibility.


Language Generation Pipeline

- Modules operate sequentially, with some dependencies.
- Each module is an independent variable.
- Enables targeted experimentation to study:
 1. Effect of phoneme inventory size
 2. Influence of phonotactic rules
 3. Impact of grammatical structure
 4. Variation in vocabulary generation

System Design and Implementation

Pipeline Overview

- Modular architecture for conlang generation.
- Core object: `LanguageDescription`.
- Pipeline managed by subclassing `Pipeline`.
- JSON-based output for module results.



- Each module subclasses `Module`.
- Implements `execute()` method.
- Can add/modify language features.
- Checks dependencies and raises errors if needed.

Phonetics Module

- Based on PHOIBLE phoneme segments.
- `MostCommonPhonemesModule`.

Phonotactics Module

- `BasicPhonotacticsModule` for basic C/V structures.
- `CustomPhonotacticsModule` supports specific phoneme constraints.

Generating the Phonetics

Phoneme	i	u	a	e	o	m	k	j	p	w	n
Alphabet	i	u	a	e	o	m	k	y	p	w	n

Phoneme	t	l	s	b	ŋ	g	h	d	r	f	ɲ	tʃ
Alphabet	t	l	s	b	ng	g	h	d	r	f	ñ	ch

Table 1: Output of the Most common phonemes module, and alphabet module.

Generating the Syllables

ba	be	bi	bo	bu	cha	che	chi	cho	chu
da	de	di	do	du	fa	fe	fi	fo	fu
ga	ge	gi	go	gu	ha	he	hi	ho	hu
ka	ke	ki	ko	ku	la	le	li	lo	lu
ma	me	mi	mo	mu	na	ne	nga	nge	ngi
ngo	ngu	ni	no	nu	pa	pe	pi	po	pu
ra	re	ri	ro	ru	sa	se	si	so	su
ta	te	ti	to	tu	wa	we	wi	wo	wu
ya	ye	yi	yo	yu	ña	ñe	ñi	ño	ñu

Table 2: Output of phonotactics (with CV Structure) and syllable builder modules

- Features via GrammaticalFeatures class.
- Uses AbstractFeature and AbstractPOS.
- Modules:
 - BaselineAgglutinativeGrammarModule
 - BaselineIsolatingGrammarModule

Defining the Grammar

- The language has the following standard word order: SOV
- The language is primarily agglutinative, i.e. complex words are formed by stringing together morphemes.
- All nouns have the following features:
 - The nominative case is not marked.
 - The accusative case is marked by the suffix: `ha`
 - ...
- Articles have the following features:
 - Definite (the) article is expressed using particle after the noun:
`yu`
 - ...
- ...
- There are no separate Yes or No words. Yes/No questions can be answered by echoing the verb in the question with or without negation.

Vocabulary Modules

- Data in VocabDictionary.
- Generation:
 - FromSourceVocabularyModule
 - ClusterTwoLevelVocabularyModule
 - ClusterAndSimplifyVocabularyModule
 - ApproximatingVocabularyModule
- Mapping:
 - RandomMappingModule

Generating the Vocabulary

ID	Word	Definition
turn.v.07	nibo	pass into a condition gradually, take on a specific property or attribute; become
doctor.n.01	ngosa	a licensed medical practitioner
family.n.01	wowu	a social unit living together
century.n.01	fako	a period of 100 years
⋮	⋮	⋮

Table 3: Output of Vocabulary Generation and Mapping Modules

- Translates using LLMs.
- Paragraph-level translation of `AbstractSourceText`.

Source

People have been painting pictures for at least 30,000 years. The earliest pictures were painted by people who hunted animals. They used to paint pictures of the animals they wanted to catch and kill. Pictures of this kind have been found on the walls of caves in France and Spain. No one knows why they were painted there. Perhaps the painters thought that their pictures would help them to catch these animals. Or perhaps human beings have always wanted to tell stories in pictures.


Translation

Yoche pado-re pula gu pofenu. Tu pula li-pado ta yoche lei necha yofe. Lei pado-re pula wi yofe lei ngegu pado da necha. Pula wi po ngono ngiñatu-re ta fuhuyu wi po ñayo France da Spain. Yo lei gasu la lei li-pado lu. Howome yu yoche huchi pa lei pula mehi lei pado po yofe. Do howome yoche ngegu-re ruba dime ta pula.

Figure 1: Output of the translation module

Evaluation Framework

- Evaluations handled by Evaluator class.
- Independent of the main pipeline.
- Outputs saved in an evaluations folder.



- Evaluation modules benchmark generated languages.
- Results stored alongside each language.
- Three main categories:
 - Information Loss
 - Simplicity
 - Zipf's Law Compliance

- **Machine Translation Scores:**
 - Round-trip translation using LLM
 - Evaluated with BLEU, ROUGE, METEOR
- **Reading Comprehension:**
 - LLM answers questions based on original and translated texts
 - Score = % of correct answers

- **BERT Fine-tuning:** Measures perplexity of model trained on new language.
- **Vocabulary Size:** Total lexicon count.
- **Phonemic Inventory Size:** Number of phonemes used.
- **Phonotactic Complexity:** Assesses rule complexity.

Zipf's Law Metric

- Evaluates how closely the generated language follows Zipf's Law.
- Zipf exponent compared against that of the original English text.
- Indicator of natural distribution of word frequencies.


- DetranslationEvaluator (BLEU, METEOR, ROUGE)
- GenerateTranslationSummary (Vocab Size, Syllables)
- CompressionEvaluator
- RaceCEvaluator
- BertEvaluator (Perplexity, Eval Loss)
- ZipfsLawEvaluator

- **Baseline:** Most common phonemes, CV rules.
- **Approximation:** Fixed vocabulary + embedding-based matching.
- **Two-Level:** Bi-syllabic clustering-based vocabulary.
- **Cluster & Simplify:** One word per cluster.
- **Phonotactics:** Custom syllable constraints.

Results and Evaluation


Effect of Phoneme Count

- Minimal impact on translation or Race-C scores.
- Suggests that languages can be simplified in phoneme inventory.
- Meaning preservation not significantly affected.




Effect of Phonotactics

- Simplified phonotactic rules do not degrade performance.
- Translation and comprehension scores largely unaffected.
- Simplification viable for phonological structure.



Effect of Grammar Rules

- Compared Agglutinative vs. Isolating grammars.
- No major difference in translation or comprehension.
- Only Bert perplexity varied—needs further analysis.



Vocabulary Simplification and Performance

- Vocabulary simplification causes performance drop.
- MTS (n-gram) suffers more than Race-C (meaning-based).
- Ratio of score to vocab size is better for simpler vocab methods.

Effect of Simplified Vocabulary




Figure 2: Race-C Accuracy





Figure 3: Mean Translation Score (MTS)

Effect of Language Model Temperature

- High temperature leads to less consistent results.
- Low to moderate values (e.g., 0.2) are stable.
- Reproducibility improves with lower temperature.



- Summary of Approximating Pipeline 18 5
- Detranslation of Cluster Simplify 90 Clusters

Conclusion

Conclusion

- LLMs can aid in building efficient, minimal conlangs.
- Simplification is possible in phonology, grammar, and vocab.
- Reading comprehension is robust even with simplification.
- Zipf-like distributions emerged naturally.

- Explore more grammar/vocab generation strategies.
- Better metrics for simplicity and efficiency.
- Deeper analysis of grammatical influence on meaning retention.
- Extend to spoken or multi-modal languages.

Questions?

Thank you!
Questions are welcome.

Comparing Simplified Vocabularies




Figure 4: Comparing different Simplifying vocabulary generation methods with the baseline.

LLMs and Embedding Models

- **Local LLMs:** via Ollama, e.g. deepseek-r1:14b.
- **Remote LLMs:** OpenAI and Groq (gpt-4o).
- **Embeddings:** all-MiniLM-L6-v2.

- Subclasses of `ClusteringMethod`.
- Vocabulary clustering: `AgglomerativeClustering`.
- Alternatives supported: `KMeans`, `DBScan`, `HDBScan`.

- Dataset: RaceC (Race-C dataset).
- Word frequency: `wordfreq` library.
- Corpus source: Exquisite Corpus.
- Outputs: HTML summaries and translation pairs.